


## HW2 Writeup

### Problem 1.2



### Problem 1.3

The smallest unit of thread is warp which consists of 32 processors and the cache line of each warp has a width of 128 bytes. Therefore, in the case of *uint*, the warp will use up all the 128 bytes loaded into the cache since the size of *uint* is 4 bytes and  $32 \times 4$  bytes = 128 bytes. However, for *char*, since each *char* variable has only a size of 1 byte, 3/4 of the data loaded into the cache is wasted in one execution of the warp. This explains why using *uint* format has higher bandwidth. However, the increase in bandwidth is lower than the expected value of 4 because the excessive data loaded into the cache in the case of *char* can be used in the next warp execution.

There is not much difference in bandwidth for both the *uint* and *uint2* cases because in both cases, there is no waste in memory. Any data loaded into the cache during each warp execution will all be used within that execution.

### Problem 2.2


In the cuda function *device\_graph\_propagate()*, we can see that for each entry output of vector *graph\_nodes\_out*, there is one write operation.

As for the read operations, it can be seen that for each node, we need to read the range of neighbors *graph\_indices[i]* and *graph\_indices[i + 1]*. For each neighbor, we also have three read operators to get *graph\_edges*, *inv\_edges\_per\_node* and *graph\_nodes\_in*. Therefore, for each node, there are  $2 + 3*avg\_edges$  read operations.

Since both data types *uint* and *float* have 4 bytes, the total number of read and write operations are:

$$(3 + 3*avg\_edges) * num\_nodes * num\_iterations * (4 \text{ bytes})$$

### Problem 2.3



### Problem 2.4

The memory access in this problem is random due to the data structure *graph\_indices* and *graph\_edges*. Once we get the indices of neighboring nodes from *graph\_edges*, we have to access the input vector *graph\_nodes\_in* and *inv\_edges\_per\_node* randomly to calculate the output value. This data access pattern is bad since we cannot reuse data loaded into cache. As we can see, the maximum bandwidth in problem 1 is

over 100 GB/s and even the smallest one is still above 20 GB/s. However, in problem 2, the highest bandwidth is only around 9 GB/s. In problem 1, contiguous memory is accessed when elements of arrays are read. This data accessing pattern is much better than the random data accessing pattern in problem 2 in caching and it explains why bandwidth in problem 1 is much higher than that in problem 2.